

Subgroups of free groups and primitive elements

A. Clifford and R. Z. Goldstein

(Communicated by J. Howie)

Abstract. An algorithm is developed to determine whether a subgroup of a free group contains a primitive element. This answers question 39b in the list of problems on free groups posted on the World of Groups website, www.grouptheory.info.

1 Introduction

The World of Groups website [12] maintains a list of open problems in group theory. Question F39 has two parts. The first part asks whether there is an algorithm which, when given a subgroup H of a free group F and an element $w \in F$, can determine whether there is an automorphism $\phi : F \rightarrow F$ so that $\phi(w) \in H$. In this paper, we address the second part by developing an algorithm which, given a subgroup H of a free group F , determines whether H contains an element which is part of some basis for F . Such an element is called *primitive*. Question F39a still remains unanswered in general. However, in a recent preprint [9], Silva and Weil have given an algorithm in the case where the ambient free group has rank two.

2 Automorphisms

Throughout, F will denote a free group on n generators. We will take as a basis for F the set $X = \{x_1, x_2, \dots, x_n\}$. We begin by discussing various automorphisms of F . Let σ be a permutation of the set $\{1, 2, \dots, n\}$ and for each i , let \hat{x}_i be a choice of either x_i or x_i^{-1} . Then we define the automorphism $\phi_{\sigma, \wedge} : F \rightarrow F$ on the basis by $\phi_{\sigma, \wedge}(x_i) = \hat{x}_{\sigma(i)}$. This is a *type I automorphism*. We note that the set of type I automorphisms is closed under composition.

Let k and l be two integers satisfying $1 \leq k \leq l \leq n$. Then a *type II automorphism* is of the form $\phi_{k,l} : F \rightarrow F$ defined on the basis $X = \{x_1, x_2, \dots, x_n\}$ by:

$$\begin{aligned}\phi_{k,l}(x_1) &= x_1, \\ \phi_{k,l}(x_i) &= x_1 x_i \text{ for } 2 \leq i \leq k, \\ \phi_{k,l}(x_i) &= x_1 x_i x_1^{-1} \text{ for } k < i \leq l, \\ \phi_{k,l}(x_i) &= x_i \text{ fixed for } l < i \leq n.\end{aligned}$$

To ensure that $\phi_{k,l}$ is not the identity, we make the further assumption that $l \geq 2$. One can easily obtain all Nielsen automorphisms as compositions of Type I and Type II automorphisms. It is well known (cf. [5]) that every automorphism of F can be written as a composition of Nielsen automorphisms. Therefore, the collection of Type I and Type II automorphisms generates the automorphism group of F .

3 Graphs

All graphs considered will be finite, directed and labeled with the elements of $X = \{x_1, x_2, \dots, x_n\}$. All graph maps between such graphs will be assumed to preserve the direction and labeling of the edges. Given a graph K , we shall use E_K to denote the edge set of K and $|K|$ to denote the cardinality of E_K .

Let K be such a graph. A path in K corresponds to an element of F in the usual way. We pick a vertex v_0 in K , and define the subgroup $H = \pi K$ as the set of elements obtained from closed paths at v_0 . If we choose a different vertex, we get a subgroup H_1 conjugate to H . Since H contains a primitive element if and only if any conjugate of it does, we will not make reference to base points and consider K to represent any subgroup in the conjugacy class of πK .

The graph K is reduced if there is no vertex v adjacent to two edges both labeled with some x_i , either both directed toward v , or both directed away from v . If a graph K is not reduced, we can perform a sequence of Stallings foldings, as described in [7], to obtain a reduced graph K^* so that $\pi K = \pi K^*$. Moreover, there is an induced graph map $p_K : K \rightarrow K^*$ which we call a folding map. Throughout, we will omit subscripts when they become too burdensome. Note that if v and w are vertices in K , then $p_K(v) = p_K(w)$ if and only if there is a path in K from v to w whose word freely reduces to 1 in F . Let $f : L \rightarrow K$ be a graph map between the graphs L and K . There is an induced map $f^* : L^* \rightarrow K^*$ so that $p_K \circ f = f^* \circ p_L$.

For any finitely generated subgroup H of F , there is a finite reduced graph K so that $H = \pi K$. To see this, let L be the covering of the wedge of n circles corresponding to H . Then we can take K to be any subgraph of L , appropriately directed and labeled, so that the inclusion induced $i_* : \pi_1(K) \rightarrow \pi_1(L)$ is an isomorphism. Conversely, if $\pi K = H$, we can recognize it as a subgraph of L . We call such a graph a *core* for H . We note that typically core graphs are precluded from having degree one vertices. However, for technical reasons, we shall find it convenient to allow core graphs to have degree one vertices as may result from our definition.

Let K be a core for H , a subgroup of F . For a given automorphism ϕ of F , we wish to describe a core for $\phi[H]$. If ϕ is a type I automorphism, clearly by relabeling some edges and by redirecting some edges, we can obtain a graph K' so that $\phi_{\sigma, \wedge}(\pi K) = \pi K'$.

We wish to define a similar process for the type II automorphism $\phi_{k,l}$. Let K be a core graph for the subgroup H as described above. For each edge e labeled x_i for some $i \in \{2, \dots, k\}$, we introduce a vertex, v_e and two edges, one from $\iota(e)$ to v_e labeled x_1 , and the other from v_e to $\tau(e)$ labeled x_i . We then remove the original edge e . Similarly, for each edge e labeled x_i for some $i \in \{k+1, \dots, l\}$, we introduce two vertices, u_e and w_e and three edges, two labeled x_1 , one from $\iota(e)$ to u_e and one from $\tau(e)$

to w_e , and one from u_e to w_e labeled x_i . We then remove the original edge e . In this process, the edge that we remove has the same label as exactly one of the two or three edges that we added. We call the added edge with the same label as the removed edge a *replacement* edge. The other added edges are all labeled x_1 . We call these edges *new* x_1 edges. The edges of K labeled x_1 are *old* x_1 edges. We call this new graph \hat{K} . Now, there is probably not a graph map from K to \hat{K} , but there is an injection $\rho_K : E_K \rightarrow E_{\hat{K}}$ where $\rho_K(e)$ is the replacement edge of e if e is labeled x_i for some i with $2 \leq i \leq l$, and $\rho_K(e) = e$ otherwise. The proof of the following lemma will be left to the reader.

Lemma 1. *Let $f : L \rightarrow K$ be a graph map. Let \hat{L} and \hat{K} be obtained from L and K via the automorphism $\phi_{k,l}$. Then there is an induced graph map $\hat{f} : \hat{L} \rightarrow \hat{K}$ so that for any edge e of L , $\hat{f}[\rho_L(e)] = \rho_K(f[e])$.*

The graph \hat{K} may not be a core subgraph as it may not be reduced. We perform a sequence of Stallings foldings to obtain a core graph $K' = (\hat{K})^*$ which may contain degree one vertices. We will find the following lemmas useful. The proof of the first is left to the reader.

Lemma 2. *Let $F[x_1, x_2, \dots, x_n] = F$ be the free group on n generators. Let H be a subgroup of F with core graph K . Let $\phi : F \rightarrow F$ be an automorphism which acts as a type II basis change. Let $K' = (\hat{K})^*$ be the core graph obtained from K via ϕ as described above. Then K' is a core graph of $\phi(H)$.*

Lemma 3. *Let K be a core graph. Let \hat{K} be obtained from K via the type II automorphism $\phi_{k,l}$ and let $p_{\hat{K}} : \hat{K} \rightarrow K'$ be a folding map. If e_1 and e_2 are two edges so that $p_{\hat{K}}(e_1) = p_{\hat{K}}(e_2)$, then both edges are labeled x_1 and at least one is a new x_1 edge.*

Proof. Let e_1 and e_2 be two edges of \hat{K} adjacent to v_1 and v_2 respectively so that $p_{\hat{K}}(e_1) = p_{\hat{K}}(e_2)$ and $p_{\hat{K}}(v_1) = p_{\hat{K}}(v_2)$. Clearly, this implies that e_1 and e_2 have the same label. There is a path in \hat{K} from v_1 to v_2 whose label is a word freely equal to 1. If both e_1 and e_2 correspond to edges of K , then there is a path between the corresponding vertices in K whose label is a word freely equal to 1. This is impossible as K is reduced. The result follows. \square

Lemma 4. *Let $f : L \rightarrow K$ be a map between reduced graphs. Let L' and K' be the reduced graphs obtained from L and K via $\phi_{k,l}$ and let $f' : L' \rightarrow K'$ be the induced graph map. If e_1 and e_2 are distinct edges of L which are identified by f , then the corresponding edges e'_1 and e'_2 are distinct and identified by f' .*

Proof. That e'_1 and e'_2 are distinct is a consequence of Lemma 3. It is a consequence of Lemma 1 and the fact that the folding maps commute with induced maps that f' identifies e'_1 and e'_2 . \square

4 Manifolds

In [10] and [11], Whitehead developed techniques for studying automorphisms of free groups using $M = (S^1 \times S^2) \# (S^1 \times S^2) \# \cdots \# (S^1 \times S^2)$, the 3-manifold which is the connected sum of n copies of $S^1 \times S^2$. We will provide a brief review of these techniques, but refer the interested reader to [3], [4] and [8] for more details. A sphere basis for M is an ordered n -tuple of spheres $\Sigma = (S_1, S_2, \dots, S_n)$ embedded in M so that $M - \bigcup_{i=1}^n S_i$ is connected. If we split each S_i , we get a *pre-manifold* M_1 which is homeomorphic to the 3-sphere S^3 with $2n$ copies of the open 2-ball B^2 removed. So each S_i of M corresponds to two boundary components of M_1 which we call S_i^- and S_i^+ . Now, M is obtained from M_1 by identifying each S_i^- to S_i^+ by reversing orientation. For clerical expedience, we let $\psi : M_1 \rightarrow M$ be that identification map. Now, let U be a collar on S_i in M . Then $U - S_i$ has two components U_1 and U_2 . If the closure of $\psi^{-1}[U_j]$ contains S^+ , we call U_j the *positive side* of S_i . Otherwise, the closure of $\psi^{-1}[U_j]$ contains S^- and we call U_j the *negative side* of S_i .

Clearly, $\pi_1(M) = F$ is isomorphic to a free group on n generators. Since we are looking for primitive elements of F , we will only be interested in elements of F up to conjugacy. For this reason, we will not pay attention to the base point of M . Each choice of sphere basis for M corresponds to a choice of generators of F as follows: Let C be a loop in M which meets exactly one element S_i of the sphere basis. We can pull C back to M_1 where it appears as an arc between S_i^+ and S_i^- . If the arc is directed from S_i^+ to S_i^- , then the curve C represents the primitive element x_i of F ; if it is directed the other way, it represents x_i^{-1} . In this way we can determine the conjugacy class of elements of F represented by a curve embedded in M .

Again, let $\Sigma = (S_1, S_2, \dots, S_m)$ be a sphere basis for M . Reordering the elements of this n -tuple, or exchanging the roles of S_j^- and S_j^+ for some j , is tantamount to choosing a different basis of $\pi_1(M) = F$. Such a change of basis corresponds to a type I automorphism.

Next, we discuss Whitehead moves on the sphere basis for M . These moves correspond to type II automorphisms.

Let S_* be a sphere properly embedded in M disjoint from the sphere basis. We will find it easier to refer to M_1 , in which we denote $\psi^{-1}(S_*)$ also as S_* . We note that S_* disconnects M_1 into two submanifolds, and assume that there is some fixed i_0 so that S_* separates $S_{i_0}^-$ from $S_{i_0}^+$. We call the submanifold containing $S_{i_0}^-$ the *inside* of S_* and the one containing $S_{i_0}^+$ the *outside* of S_* . By a type I basis change, we can assume that $i_0 = 1$ and that there are k and l so that the following assertions hold: for all $i \in \{1, \dots, k\}$, S_i^- is in the inside of S_* and S_i^+ is in the outside of S_* ; for all $i \in \{k+1, \dots, l\}$, both S_i^- and S_i^+ are in the inside of S_* ; and for all $i \in \{l+1, \dots, n\}$, both S_i^- and S_i^+ are in the outside of S_* .

We may use S_* to make a type II basis change by replacing S_1 with S_* . We begin by splitting M_1 along S_* to obtain a disconnected manifold D with $2n+2$ boundary components. This manifold has two components, one which contains S_1^- and a copy of S_* which we call S_*^+ ; and one which contains S_1^+ and a copy of S_* which we call S_*^- . We identify S_1^- to S_1^+ using ψ to obtain M_2 . We point out that by identifying S_*^- to S_*^+ in the obvious way and identifying S_i^- to S_i^+ using ψ , we regain the manifold

M , now with sphere basis $\Sigma' = (S_*, S_2, S_3, \dots, S_n)$. With this sphere basis, a curve that only meets S_* corresponds either to x_1 or x_1^{-1} , depending on how it is directed. So this process induces the automorphism $\phi_{k,l}$ on F .

5 Curves in M

Let $\alpha : S^1 \rightarrow M$ be an embedding of a simple closed curve C in M which is transverse to the elements of the sphere basis $\Sigma = (S_1, S_2, S_3, \dots, S_n)$ representing the cyclically reduced element $w = y_1 y_2 \dots y_m$ of F written in the basis $X = \{x_1, x_2, \dots, x_n\}$, where $y_j \in X \cup X^{-1}$. We will impose a graph structure on C using Σ . We will call the resulting graph C_Σ . An *arc of C relative to Σ* is the closure of a component of $C - \Sigma$. In terms of the identification map $\psi : M_1 \rightarrow M$, an arc is a component of $\psi^{-1}[C]$. The ordering on S^1 gives us an ordering of the arcs C_1, C_2, \dots, C_m of C , and directs each arc. We choose a point q_j on C_j . These points will be the vertices of C_Σ .

We call the components of $C - \{q_1, q_2, \dots, q_m\}$ *intervals*. The intervals of C correspond to the edges of C_Σ in the following way. The interval I_j of C is that part of C bounded by q_j and q_{j+1} . It meets exactly one element S_i from the sphere basis where $y_j = x_i^{\pm 1}$. The corresponding edge of C_Σ is labeled x_i and is directed from the vertex in the arc adjacent to S_i^- to the vertex in the arc adjacent to S_i^+ . We note that the graph C_Σ is a circle graph.

Let K be a core for the subgroup H and assume that $w \in H$. Then there is a graph map $\gamma : C_\Sigma \rightarrow K$ which reads w . The subgraph $\gamma[C_\Sigma]$ of K is called *the support of C_Σ in K* . The number of edges in the support of C_Σ in K is the *width of C_Σ in K* , denoted $w_K(C)$. We define the *deficiency of C_Σ in K* to be $\delta(C_\Sigma) = |C_\Sigma| - w_K(C)$. The deficiency measures the number of repeated edges obtained when reading the word w in K . In particular, if $\delta(C_\Sigma) = 0$, then the word w has no repeated edges in K .

Clearly the structure that we have put on the graph C_Σ depends on our choice of sphere basis $\Sigma = (S_1, S_2, S_3, \dots, S_n)$. Next we will describe how this structure changes when we choose a new sphere basis via a Whitehead move. To this end, let S_* be a sphere embedded in M disjoint from the elements of Σ , and let $\Sigma' = (S_*, S_2, S_3, \dots, S_n)$ be a sphere basis so that the change of basis from Σ to Σ' corresponds to the automorphism $\phi_{k,l}$. We assume that the arcs C_j are transverse to S_* and that no arc intersects S_* more than once. While this latter assumption is quite strong, it will be the case in our eventual application. It ensures that the word described by C is reduced with respect to the new sphere basis $\Sigma' = (S_*, S_2, \dots, S_n)$. We also will make the inconsequential assumption that no vertex of an arc lies on S_* .

The next lemma follows directly from Lemma 3.

Lemma 5. *Let C_Σ and $C_{\Sigma'}$ be the circle graphs corresponding to the curve C in M with sphere bases $\Sigma = (S_1, S_2, \dots, S_n)$ and $\Sigma' = (S_*, S_2, \dots, S_n)$ respectively where Σ' is obtained from Σ by a Whitehead move corresponding to the automorphism $\phi_{k,l}$ of F . Let $(C_\Sigma)' = (\hat{C}_\Sigma)^*$ be the graph obtained by applying $\phi_{k,l}$ to C_Σ as described in Section 3. Then $(C_\Sigma)'$ is precisely the circle graph $C_{\Sigma'}$ together with some edges labeled x_1 each adjacent to a vertex of degree one. In this way, we can consider $C_{\Sigma'}$ as a subgraph of $(C_\Sigma)'$.*

Given a sphere basis $\Sigma = (S_1, S_2, \dots, S_n)$ and a curve C_Σ , we say the sphere S_* is *stingy* with respect to C_Σ if every arc of C_Σ that meets S_* also meets S_1^- . For the rest of this paper, we will assume that our sphere S_* is stingy.

Next, assume that C represents an element of some subgroup H of F and let K be a core for H . In this situation, we have two graph maps $\gamma: C_\Sigma \rightarrow K$ and $\gamma': C_{\Sigma'} \rightarrow K'$. We wish to compare the width of C in K to its width in K' .

Theorem 1. *Let C be a curve embedded in M and let C_Σ be the circle graph corresponding to C with respect to the sphere basis $\Sigma = (S_1, S_2, \dots, S_n)$. Let S_* be a stingy sphere embedded in M so that the change of basis from Σ to $\Sigma' = (S_*, S_2, \dots, S_n)$ induces the automorphism $\phi_{k,l}$. Further, assume that there is a subgroup H with core graph K so that C_Σ represents an element $w \in H$. Let K' be the core graph obtained from K via $\phi_{k,l}$, and let $C_{\Sigma'}$ be the circle graph corresponding to the curve C with respect to the basis Σ' . Then $w_{K'}(C) \leq w_K(C)$.*

Proof. From the above lemma, we know that $(C_\Sigma)'$ is precisely the circle graph $C_{\Sigma'}$ together with edges labeled x_1 attached to some of the vertices of $C_{\Sigma'}$. Since S_* is stingy, it follows that in $C_{\Sigma'}$ every new x_1 edge is adjacent to either another new x_1 edge or an old x_1 edge with the opposite orientation. So every new x_1 edge gets folded. Therefore, $|C_{\Sigma'}| \leq |C_\Sigma|$. Now from Lemma 4, we know that if γ identifies e_1 and e_2 for some $e_1, e_2 \in E_{C_\Sigma}$ then γ' identifies e'_1 and e'_2 . The result follows. \square

6 Primitive curves in M

A curve C in M is *primitive* if it represents a primitive element of F . The following theorem will help us detect primitive curves in M . A slightly incorrect version of this statement was given in [1]. We include a corrected version here.

Theorem 2. *Let w be a reduced element of F . Then w is primitive if and only if there is a primitive curve C in M representing w with respect to the basis Σ , and an embedding of a sphere T so that $C \cap T$ consists of a single point.*

Proof. If w is primitive then there exists an automorphism $\phi: F \rightarrow F$ such that $\phi(x_1) = w$. Moreover, there is a homeomorphism $f: M \rightarrow M$ that induces ϕ . Whitehead in [11] has shown that we can take f so that $T = f[S_1]$ is normal to Σ . This means that in M_1 , no component of $\psi^{-1}[T - \Sigma]$ has a closure that meets a boundary sphere in more than one circle. Let C_1 be a simple closed curve representing x_1 which intersects the element S_1 of the sphere basis Σ exactly once and does not meet any other element of Σ . We note that $f[C_1]$ meets T exactly once and the word which $f[C_1]$ represents is equivalent to w in F although it may not be reduced. Now choose a pair C and T satisfying the following conditions: the word which C represents is equal to w in F ; T is normal to Σ ; $C \cap T$ has exactly one point; and C has the minimal number of intersections with the sphere basis Σ among such pairs. We claim that the word that C represents is reduced.

To see this, assume that it is not reduced. Then there is a subarc γ of C with end points p and q so that p and q are on some S_i ; no other point of γ is on an element of Σ ; and in M_1 the closure of $\psi^{-1}[\gamma]$ has its two boundary points on the same boundary sphere of M_1 , either S_i^- or S_i^+ . Without loss of generality, we will assume that it is the latter.

Now, the sphere T meets S_i in a collection of circles $\{D_1, D_2, \dots, D_r\}$. We will assume that neither p nor q is on any of these circles. If p and q are on the same component of $S_i - \bigcup D_j$, then γ does not meet T because $\gamma \cup \gamma'$, being trivial, must meet T in an even number of points. So we can obtain a primitive curve C' representing w meeting T in one point with fewer intersections with Σ as follows: let γ' be a path from p to q on S_i which is disjoint from T . Now, let $C' = (C - \gamma) \cup \gamma'$ and alter C' by a homotopy to push γ' into the negative side of S_i . This contradicts the minimality of C .

Next, assume there is some D_j separating p and q on S_i . In this case, since T is normal to Σ , there is a component Δ of $T - \Sigma$ whose closure contains D_j so that γ meets Δ in a point y . Clearly, y must be the unique point in $T \cap C$. We see that the normality of T also implies that if there is such a D_j , there is only one. Here, let γ' be a path on S_i from p to q that meets D_j in one point and no other points of T . As above, let $C' = (C - \gamma) \cup \gamma'$ and alter C' by a homotopy to push γ' into the negative side of S_i . Again we see that C' meets T in one point and represents the word w in F . Moreover, C' has fewer points of intersection with Σ than does C .

Now, to prove the converse, suppose that there is a curve C which represents w with respect to Σ and a 2-sphere T such that C intersects T in precisely one point. According to [3], T is a primitive sphere, since the complement of T in M is connected. Thus we can find 2-spheres $\Sigma_2, \dots, \Sigma_n$ so that $\Sigma' = \{T, \Sigma_2, \dots, \Sigma_n\}$ is a sphere basis. Thus the sphere bases Σ and Σ' represent an automorphism ϕ of F . Then label the sphere T as x_1 and label the other spheres accordingly. Now $\phi(w)$ has exactly one occurrence of x_1 or exactly one occurrence of x_1^{-1} but not both. Thus w is primitive. \square

Such a sphere T as described in this theorem is called a *detection sphere* for C . Let $\Sigma = (S_1, S_2, \dots, S_n)$ be a sphere basis of M . We assume that any detection sphere is transverse to Σ . It is not true that every primitive curve has a detection sphere. If the primitive curve C has a detection sphere, then we say that C is a *proper primitive curve*. If C is a proper primitive curve, the *complexity of C relative to Σ* , denoted $\chi_\Sigma(C)$, is the minimal number of components of $T \cap \bigcup S_i$ among detection spheres T for C .

Theorem 3. *Let K be a core graph of some subgroup H . Let C be a proper primitive curve in M representing the primitive element w of H , and let $\Sigma = (S_1, S_2, \dots, S_n)$ be a sphere basis. Assume that $\chi_\Sigma(C) > 0$. Then there is another sphere basis for M , Σ' obtained from Σ by a composition ϕ of a type I automorphism and a type II automorphism so that $\chi_{\Sigma'}(C) < \chi_\Sigma(C)$ and $w_{K'}(C) \leq w_K(C)$, where K' is the graph obtained from K via ϕ .*

Proof. Let T be the detection sphere realizing the complexity of C . Then, since $\chi_\Sigma(C) > 0$, T must meet the sphere basis Σ . An endcap E is a subset of T which is homeomorphic to a closed disk so that the boundary of E lies on some S_i and no point on the interior of E lies on any element of Σ . Clearly, T has at least two endcaps, at least one of which does not meet C . Let us denote this endcap as E .

We consider E as a subset of M_1 . Then the boundary of E meets either S_i^+ or S_i^- . Without loss of generality, we assume it is the latter. Moreover, E separates some elements of Σ^\pm from the others, since otherwise E would be a trivial endcap as described in [10], in which case we could alter T by an isotopy to reduce $\chi_\Sigma(C)$.

Let $j: S^2 \times I \rightarrow M_1$ be an embedding describing a collar of S_i^- so that $j[S^2 \times \{0\}] = S_i^-$. We let J be the image of j and let J_1 be $j[S^2 \times \{1\}]$. We take this collar small enough so that the following conditions hold: C and T are transverse to J_1 ; $E \cap J_1$ has one component; $T \cap J$ is a collection of annuli each of which has one boundary component on each of S_i^- and J_1 ; and $C \cap J$ consists of a collection of subarcs of C each of which meet S_i^- .

Now E meets J_1 in a circle. We let Δ be the disk on J_1 bounded by this circle so that $E \cup \Delta$ separates S_i^+ from S_i^- . We let S_* be the sphere $(E - J) \cup \Delta$. Clearly, S_* is not transverse to T , but we can push S_* slightly so that S_* is disjoint from E . Now S_* is a sphere embedded in M_1 and every arc of C that meets S_* also meets S_i^- .

Using a type I automorphism, we can assume that $i = 1$ and that there are k and l so that the following assertions hold: for all $i \in \{1, \dots, k\}$, S_i^- is in the inside of S_* and S_i^+ is in the outside of S_* ; for all $i \in \{k+1, \dots, l\}$, both S_i^- and S_i^+ are in the inside of S_* ; and for all $i \in \{l+1, \dots, n\}$, both S_i^- and S_i^+ are in the outside of S_* . We note that S_* is stingy.

Now we do a Whitehead move corresponding to $\phi_{k,l}$ to obtain the basis $\Sigma' = (S_*, S_2, \dots, S_n)$ and the graph K' . Since S_* does not meet E , we have $\chi_{\Sigma'}(C) < \chi_\Sigma(C)$. Moreover, by Theorem 1, we know that $w_{K'}(C) \leq w_K(C)$. \square

Theorem 4. *Let K be a core graph of some subgroup H . Let C be a proper primitive curve in M representing the primitive element w of H , and let $\Sigma = (S_1, S_2, \dots, S_n)$ be a sphere basis. Assume that $\chi_\Sigma(C) = 0$. Then there is another proper primitive curve C_1 so that $\chi_\Sigma(C_1) = 0$ and the deficiency $\delta_K(C_1) = 0$.*

Proof. We choose a proper primitive curve C_1 so that $\chi_\Sigma(C_1) = 0$ and $\delta_K(C_1)$ is as small as possible. By way of contradiction, assume $\delta_K(C_1) > 0$. Then there are two intervals I_1 and I_2 of C_1 so that $\gamma(I_1) = \gamma(I_2)$. We assume that this edge of K is labeled x_i so that there are points $p_1 \in S_i \cap I_1$ and $p_2 \in S_i \cap I_2$. These points separate C_1 into two intervals D_1 and D_2 .

Since C_1 is a proper primitive curve and $\chi_\Sigma(C_1) = 0$, there is some detection sphere T so that T is disjoint from each element of Σ and T meets C_1 in exactly one point, which, without loss of generality, we assume is on D_1 . Since T does not meet S_i , there is a small interval D_3 on S_i from p_1 to p_2 which is disjoint from T . We define the curve $C_2 = D_1 \cup D_3$ and alter C_2 by a homotopy near D_3 so that it is transverse to S_i . Now C_2 meets T in exactly one point, so it is a proper primitive curve. Moreover, the construction ensures that there is a corresponding map $\gamma_2: C_2 \rightarrow K$. Lastly,

we note that $\delta_K(C_2) < \delta_K(C_1)$. This contradicts the minimality of C_1 . Therefore, $\delta_K(C_1) = 0$. \square

7 The algorithm

Let $X = \{x_1, x_2, \dots, x_n\}$ be a basis for the free group F_n . In [10], Whitehead gave an algorithm which determines whether or not a word w in the alphabet X is a primitive element of F_n . Significant improvements in the complexity of this algorithm have been given in [6].

Let K be a core graph of the subgroup H of F . Determining whether H contains a primitive element is equivalent to determining whether there is a path in K that represents a primitive element. Such a path will be called a *primitive path*. A path is *simple* if it uses no edge more than once. Clearly, there is an algorithm to determine all the simple paths in K . Since there is an algorithm which determines whether a given path is a primitive path, it is possible to determine whether K has a simple primitive path.

Given the graph K with N edges, a graph K' is *one step away from K* if there is a graph K'' so that the following conditions hold: K' is subgraph of K'' ; K' has no more than N edges; and K'' is obtained from K by performing a type I and a type II automorphism on K . We see that if K' is one step away from K and K' contains a primitive path, then so does K . (We point out that it is possible that K contains a primitive path while K' does not.) Clearly, given K there is an algorithm which produces all graphs which are one step away from K . We say that K' is M steps away from K if there is a sequence of graphs $K = K_0, K_1, \dots, K_M = K'$ so that each K_j is one step away from K_{j-1} and M is minimal in this regard.

Given K , let

$$\Omega_K = \{K' \mid \text{there is some } M \text{ so that } K' \text{ is } M \text{ steps away from } K\}.$$

Since each $K' \in \Omega_K$ has no more than N edges, Ω_K is a finite class of graphs. Also, if some $K' \in \Omega_K$ contains a primitive path, then so does K . Moreover, $K' \in \Omega_K$ if and only if $\Omega_{K'} \subset \Omega_K$.

Now there is an algorithm which produces all of the elements of Ω_K . So there is an algorithm which can determine whether there is an element of Ω_K that has a simple primitive path. This is the algorithm that determines whether K has a primitive path as is made clear in the next theorem.

Theorem 5. *K has a primitive path if and only if there is some K' in Ω_K that has a simple primitive path.*

Proof. Clearly, if there is some K' in Ω_K that has a simple primitive path, then K has a primitive path.

Conversely, assume that K has a primitive path α . Let C be a proper primitive curve representing α embedded in M with sphere basis Σ and detection sphere T so

that $\chi_\Sigma(C)$ is as small as possible. We proceed by induction on $\chi_\Sigma(C)$. If $\chi_\Sigma(C) = 0$, then by Theorem 3, K contains a simple primitive path.

If $\chi_\Sigma(C) > 0$, then by Theorem 2, there is another sphere basis Σ' for M , obtained from Σ by a composition ϕ of a type I automorphism and a type II automorphism so that $\chi_{\Sigma'}(C) < \chi_\Sigma(C)$ and $w_{K''}(C) \leq w_K(C)$, where K'' is the graph obtained from K via ϕ . We let K''' be the support of C in K'' . Clearly $K''' \in \Omega_K$. Since $\chi_{\Sigma'}(C) < \chi_\Sigma(C)$, there is some $K' \in \Omega_{K'''}$ that contains a simple primitive path. But $\Omega_{K'''} \subset \Omega_K$, so $K' \in \Omega_K$. \square

We end this note with a brief discussion of the complexity of this algorithm. Let K be a graph with n edges whose rank is r . If α is a simple path in K , clearly α has no more than n edges. In [6], it is shown that it can be determined whether or not α represents a primitive element in linear time in n . Unfortunately, the number of simple paths in K is of the order of $(2r)^r$. Lastly, we note that for the algorithm to conclude that K does not have a primitive path, we need to check every simple path not only in K , but in each graph in Ω_K , which although finite, can be quite large. In conclusion, we see that the algorithm described here is very time consuming. While some efficiencies in this algorithm may be obtainable, we think that if ever a reasonably practical algorithm is found, it will be considerably different.

Acknowledgement. The authors would like to thank the referee for numerous suggestions which led to significant improvements to this paper. In particular, the referee informed us about the existence of the very interesting preprint [9].

References

- [1] A. Clark and R. Goldstein. On a method of J. H. C. Whitehead. *Comm. Algebra* **34** (2006), 1615–1623.
- [2] S. M. Gersten. Reducible diagrams and equations over groups. In *Essays in group theory*, Math. Sci. Res. Inst. Publ. 8 (Springer, 1987), pp. 15–73.
- [3] R. Z. Goldstein. The length and thickness of words in a free group. *Proc. Amer. Math. Soc.* **127** (1999), 2857–2863.
- [4] R. Z. Goldstein and E. C. Turner. Monomorphisms of finitely generated free groups have finitely generated equalizers. *Invent. Math.* **82** (1985), 283–289.
- [5] R. C. Lyndon and P. E. Schupp. *Combinatorial group theory* (Springer-Verlag, 1977).
- [6] A. Roig, E. Ventura and P. Weil. On the complexity of the Whitehead minimization problem. *Internat. J. Algebra Comput.* **17** (2007), 1611–1634.
- [7] J. R. Stallings. Topology of finite graphs. *Invent. Math.* **71** (1983), 551–565.
- [8] J. R. Stallings. Whitehead graphs on handlebodies. In *Geometric group theory down under (Canberra, 1996)* (de Gruyter, 1999), pp. 317–330.
- [9] P. V. Silva and P. Weil. Automorphic orbits in free groups: words versus subgroups. Preprint (2009). <http://arxiv.org/abs/0809.4386>.
- [10] J. H. C. Whitehead. On certain sets of elements in a free group. *Proc. London Math. Soc.* (2) **41** (1936), 48–56.
- [11] J. H. C. Whitehead. On equivalent sets of elements in a free group. *Ann. of Math.* (2) **37** (1936), 782–800.
- [12] World of Groups website. <http://www.grouptheory.info/>

Received 3 February, 2009; revised 21 August, 2009

A. Clifford, Department of Mathematics and Statistics, The College of New Jersey, P.O. Box 7718, Ewing, NJ 08628-0718, U.S.A.

E-mail: cliffan@tcnj.edu

R. Z. Goldstein, Department of Mathematics, 1400 Washington Avenue, Albany, NY 12222, U.S.A.

E-mail: richardg@csc.albany.edu